

Randomized algorithms for least squares and matrix approximation

Tyler Chen

August 26, 2024

chen.pw/slides

About me

I work on the design and analysis of **practically fast** and **theoretically justified randomized algorithms** for fundamental linear algebra tasks.

I like working with nearby communities (physics, theoretical computer science, optimization, computational science, etc.).

Academic history:

- Currently an **Assistant Professor / Courant Instructor** at New York University
 - Sponsor: Chris Musco
- PhD in Applied Math at University of Washington
 - Advisors: Anne Greenbaum and Tom Trogon
- B.S. in Math and Physics at Tufts University, minor in Studio Art

Today

Stochastic Optimization¹

- First proof of $O(\sqrt{k})$ convergence of minibatch stochastic gradient descent with heavy-ball momentum

Matrix approximation/operator learning²

- First matrix-vector product algorithm for near-optimal hierarchical matrix approximation
- Matrix-vector product algorithms for approximating fixed-sparsity matrices (no dimension dependence and matching lower bounds!)

¹Bollapragada, Chen, and Ward 2024.

²Amsel, Chen, Keles, Halikias, Musco, and Musco 2024; Chen, Keles, Halikias, Musco, Musco, and Persson 2024.

Linear least squares

Consider a consistent least squares problem

$$\min_{\mathbf{x} \in \mathbb{R}^d} f(\mathbf{x}), \quad f(\mathbf{x}) = \frac{1}{2} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2 = \sum_{i=1}^n \frac{1}{2} (\mathbf{a}_i^\top \mathbf{x} - b_i)^2.$$

The gradient is

$$\nabla f(\mathbf{x}) = \mathbf{A}^\top (\mathbf{A}\mathbf{x} - \mathbf{b}) = \sum_{i=1}^n \mathbf{a}_i (\mathbf{a}_i^\top \mathbf{x} - b_i).$$

We can then implement **gradient descent**

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha_k \nabla f(\mathbf{x}_k).$$

This requires one product with each \mathbf{A} and \mathbf{A}^\top every iteration.

Improvements to gradient descent

Stochastic gradients: Often we don't want to read the whole data matrix each iteration. Instead, we sample a random index i with probability $p_i \propto \|\mathbf{a}_i\|^2$, and then use the stochastic gradient

$$\nabla f_i(\mathbf{x}) = \frac{1}{p_i} \mathbf{a}_i (\mathbf{a}_i^\top \mathbf{x} - b_i).$$

The stochastic gradient is equal to the true gradient **on average**:

$$\mathbb{E}[\nabla f_i(\mathbf{x})] = \sum_{i=1}^n p_i \cdot \frac{1}{p_i} \mathbf{a}_i (\mathbf{a}_i^\top \mathbf{x} - b_i) = \sum_{i=1}^n \mathbf{a}_i (\mathbf{a}_i^\top \mathbf{x} - b_i) = \nabla f(\mathbf{x}).$$

Momentum: Instead of just updating based on the gradient, we can take previous iterates into account.

Convergence guarantees

Both momentum and stochastic gradient improve on classical gradient descent.

algorithm	iterations	cost/iter	formula
Gradient Descent	$\lambda_{\max}/\lambda_{\min}$	nd	$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \nabla f(\mathbf{x}_k)$
Momentum	$\sqrt{\lambda_{\max}/\lambda_{\min}}$	nd	$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \nabla f(\mathbf{x}_k) + \beta(\mathbf{x}_k - \mathbf{x}_{k-1})$
Stochastic Gradient	$n \cdot (\lambda_{\text{ave}}/\lambda_{\min})$	d	$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \nabla f_i(\mathbf{x}_k)$

However, the bounds are somewhat incomparable:

- average condition number vs. square root of condition number
- the cost per iteration is not necessarily indicative of real-world costs
- sampling assumptions

Convergence guarantees

Both momentum and stochastic gradient improve on classical gradient descent.

algorithm	iterations	cost/iter	formula
Gradient Descent	$\lambda_{\max}/\lambda_{\min}$	nd	$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \nabla f(\mathbf{x}_k)$
Momentum	$\sqrt{\lambda_{\max}/\lambda_{\min}}$	nd	$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \nabla f(\mathbf{x}_k) + \beta(\mathbf{x}_k - \mathbf{x}_{k-1})$
Stochastic Gradient	$n \cdot (\lambda_{\text{ave}}/\lambda_{\min})$	d	$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \nabla f_i(\mathbf{x}_k)$

However, the bounds are somewhat incomparable:

- average condition number vs. square root of condition number
- the cost per iteration is not necessarily indicative of real-world costs
- sampling assumptions

Convergence guarantees

Both momentum and stochastic gradient improve on classical gradient descent.

algorithm	iterations	cost/iter	formula
Gradient Descent	$\lambda_{\max}/\lambda_{\min}$	nd	$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \nabla f(\mathbf{x}_k)$
Momentum	$\sqrt{\lambda_{\max}/\lambda_{\min}}$	nd	$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \nabla f(\mathbf{x}_k) + \beta(\mathbf{x}_k - \mathbf{x}_{k-1})$
Stochastic Gradient	$n \cdot (\lambda_{\text{ave}}/\lambda_{\min})$	d	$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \nabla f_i(\mathbf{x}_k)$

However, the bounds are somewhat incomparable:

- average condition number vs. square root of condition number
- the cost per iteration is not necessarily indicative of real-world costs
- sampling assumptions

Convergence guarantees

Both momentum and stochastic gradient improve on classical gradient descent.

algorithm	iterations	cost/iter	formula
Gradient Descent	$\lambda_{\max}/\lambda_{\min}$	nd	$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \nabla f(\mathbf{x}_k)$
Momentum	$\sqrt{\lambda_{\max}/\lambda_{\min}}$	nd	$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \nabla f(\mathbf{x}_k) + \beta(\mathbf{x}_k - \mathbf{x}_{k-1})$
Stochastic Gradient	$n \cdot (\lambda_{\text{ave}}/\lambda_{\min})$	d	$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \nabla f_i(\mathbf{x}_k)$

However, the bounds are somewhat incomparable:

- average condition number vs. square root of condition number
- the cost per iteration is not necessarily indicative of real-world costs
- sampling assumptions

Key question:

What is the role of momentum in SGD?

Minibatch gradient descent

In practice: momentum is used with stochastic methods and plays a critical role in training deep learning models!

In theory: momentum cannot be used to accelerate SGD.³

Minibatching: Instead of sampling a single index, we can sample several B indices at once and put them into a list S . This gives the minibatch stochastic gradient

$$\nabla f_S(\mathbf{x}) = \frac{1}{B} \sum_{i \in S} \frac{1}{p_i} \mathbf{a}_i (\mathbf{a}_i^\top \mathbf{x} - b_i).$$

One minibatch evaluation costs $O(Bd)$ operations.⁴

³Jain, Kakade, Kidambi, Netrapalli, and Sidford 2018.

⁴Due to parallelism, the real-world cost is often sort of independent of B when B is reasonably small.

Minibatch gradient descent

In practice: momentum is used with stochastic methods and plays a critical role in training deep learning models!

In theory: momentum cannot be used to accelerate SGD.³

Minibatching: Instead of sampling a single index, we can sample several B indices at once and put them into a list S . This gives the minibatch stochastic gradient

$$\nabla f_S(\mathbf{x}) = \frac{1}{B} \sum_{i \in S} \frac{1}{p_i} \mathbf{a}_i (\mathbf{a}_i^\top \mathbf{x} - b_i).$$

One minibatch evaluation costs $O(Bd)$ operations.⁴

³Jain, Kakade, Kidambi, Netrapalli, and Sidford 2018.

⁴Due to parallelism, the real-world cost is often sort of independent of B when B is reasonably small.

Minibatching+momentum

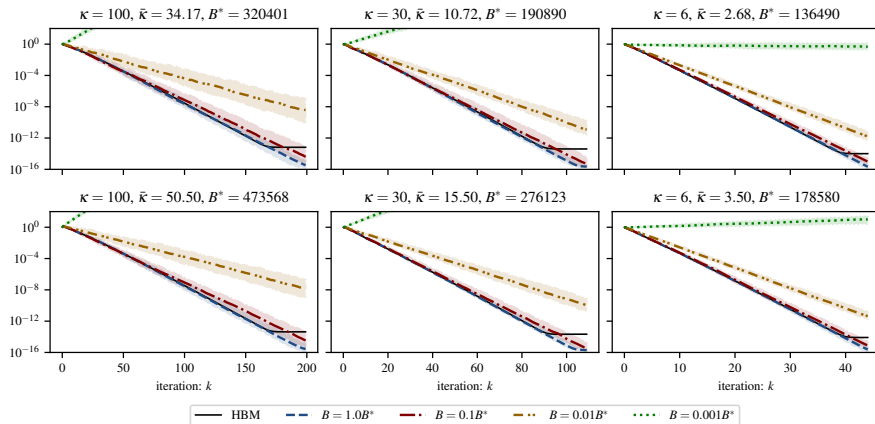
Theorem (Bollapragada, Chen, and Ward 2024). Minibatch heavy ball momentum converges in $\sqrt{\lambda_{\max}/\lambda_{\min}} \cdot \log(1/\varepsilon)$ iterations with minibatch size $B = O(d \log(d) \cdot (\lambda_{\text{ave}}/\lambda_{\min}) \cdot \sqrt{\lambda_{\max}/\lambda_{\min}})$.

Minibatch SGD can get the best of both worlds!

- same rate of convergence as full gradient methods
- cheaper gradients computations when n is large

Numerical example

Our theory might be precise enough to predict practical performance.



Proof overview (Polyak HBM)

Let's first look at how to analyze the classical heavy ball momentum algorithm:

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \underbrace{\alpha (\mathbf{A}^\top \mathbf{A} \mathbf{x}_k - \mathbf{A}^\top \mathbf{b})}_{\text{gradient}} + \underbrace{\beta (\mathbf{x}_k - \mathbf{x}_{k-1})}_{\text{momentum}}.$$

Assuming a consistent system, $\mathbf{b} = \mathbf{A} \mathbf{x}^*$, so we can rewrite the update as

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \mathbf{A}^\top \mathbf{A} (\mathbf{x}_k - \mathbf{x}^*) + \beta (\mathbf{x}_k - \mathbf{x}_{k-1}).$$

Therefore,

$$\begin{aligned} \mathbf{x}_{k+1} - \mathbf{x}^* &= (\mathbf{I} - \alpha \mathbf{A}^\top \mathbf{A}) (\mathbf{x}_k - \mathbf{x}^*) + \beta (\mathbf{x}_k - \mathbf{x}^* + \mathbf{x}^* - \mathbf{x}_{k-1}) \\ &= ((1 + \beta) \mathbf{I} - \alpha \mathbf{A}^\top \mathbf{A}) (\mathbf{x}_k - \mathbf{x}^*) - \beta (\mathbf{x}_k - \mathbf{x}_{k-1}). \end{aligned}$$

We can write this as a matrix iteration

$$\begin{bmatrix} \mathbf{x}_{k+1} - \mathbf{x}^* \\ \mathbf{x}_k - \mathbf{x}^* \end{bmatrix} = \underbrace{\begin{bmatrix} (1 + \beta) \mathbf{I} - \alpha \mathbf{A}^\top \mathbf{A} & -\beta \mathbf{I} \\ \mathbf{I} & \mathbf{0} \end{bmatrix}}_{\mathbf{T} = \mathbf{T}(\alpha, \beta)} \begin{bmatrix} \mathbf{x}_k - \mathbf{x}^* \\ \mathbf{x}_{k-1} - \mathbf{x}^* \end{bmatrix} = \mathbf{T}^k \begin{bmatrix} \mathbf{x}_1 - \mathbf{x}^* \\ \mathbf{x}_0 - \mathbf{x}^* \end{bmatrix}$$

Proof overview (Polyak HBM)

Let's first look at how to analyze the classical heavy ball momentum algorithm:

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \underbrace{\alpha (\mathbf{A}^\top \mathbf{A} \mathbf{x}_k - \mathbf{A}^\top \mathbf{b})}_{\text{gradient}} + \underbrace{\beta (\mathbf{x}_k - \mathbf{x}_{k-1})}_{\text{momentum}}.$$

Assuming a consistent system, $\mathbf{b} = \mathbf{A} \mathbf{x}^*$, so we can rewrite the update as

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \mathbf{A}^\top \mathbf{A} (\mathbf{x}_k - \mathbf{x}^*) + \beta (\mathbf{x}_k - \mathbf{x}_{k-1}).$$

Therefore,

$$\begin{aligned} \mathbf{x}_{k+1} - \mathbf{x}^* &= (\mathbf{I} - \alpha \mathbf{A}^\top \mathbf{A}) (\mathbf{x}_k - \mathbf{x}^*) + \beta (\mathbf{x}_k - \mathbf{x}^* + \mathbf{x}^* - \mathbf{x}_{k-1}) \\ &= ((1 + \beta) \mathbf{I} - \alpha \mathbf{A}^\top \mathbf{A}) (\mathbf{x}_k - \mathbf{x}^*) - \beta (\mathbf{x}_k - \mathbf{x}_{k-1}). \end{aligned}$$

We can write this as a matrix iteration

$$\begin{bmatrix} \mathbf{x}_{k+1} - \mathbf{x}^* \\ \mathbf{x}_k - \mathbf{x}^* \end{bmatrix} = \underbrace{\begin{bmatrix} (1 + \beta) \mathbf{I} - \alpha \mathbf{A}^\top \mathbf{A} & -\beta \mathbf{I} \\ \mathbf{I} & \mathbf{0} \end{bmatrix}}_{\mathbf{T} = \mathbf{T}(\alpha, \beta)} \begin{bmatrix} \mathbf{x}_k - \mathbf{x}^* \\ \mathbf{x}_{k-1} - \mathbf{x}^* \end{bmatrix} = \mathbf{T}^k \begin{bmatrix} \mathbf{x}_1 - \mathbf{x}^* \\ \mathbf{x}_0 - \mathbf{x}^* \end{bmatrix}$$

Proof overview (Polyak HBM)

Let's first look at how to analyze the classical heavy ball momentum algorithm:

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \underbrace{\alpha (\mathbf{A}^\top \mathbf{A} \mathbf{x}_k - \mathbf{A}^\top \mathbf{b})}_{\text{gradient}} + \underbrace{\beta (\mathbf{x}_k - \mathbf{x}_{k-1})}_{\text{momentum}}.$$

Assuming a consistent system, $\mathbf{b} = \mathbf{A} \mathbf{x}^*$, so we can rewrite the update as

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \mathbf{A}^\top \mathbf{A} (\mathbf{x}_k - \mathbf{x}^*) + \beta (\mathbf{x}_k - \mathbf{x}_{k-1}).$$

Therefore,

$$\begin{aligned} \mathbf{x}_{k+1} - \mathbf{x}^* &= (\mathbf{I} - \alpha \mathbf{A}^\top \mathbf{A}) (\mathbf{x}_k - \mathbf{x}^*) + \beta (\mathbf{x}_k - \mathbf{x}^* + \mathbf{x}^* - \mathbf{x}_{k-1}) \\ &= ((1 + \beta) \mathbf{I} - \alpha \mathbf{A}^\top \mathbf{A}) (\mathbf{x}_k - \mathbf{x}^*) - \beta (\mathbf{x}_k - \mathbf{x}_{k-1}). \end{aligned}$$

We can write this as a matrix iteration

$$\begin{bmatrix} \mathbf{x}_{k+1} - \mathbf{x}^* \\ \mathbf{x}_k - \mathbf{x}^* \end{bmatrix} = \underbrace{\begin{bmatrix} (1 + \beta) \mathbf{I} - \alpha \mathbf{A}^\top \mathbf{A} & -\beta \mathbf{I} \\ \mathbf{I} & \mathbf{0} \end{bmatrix}}_{\mathbf{T} = \mathbf{T}(\alpha, \beta)} \begin{bmatrix} \mathbf{x}_k - \mathbf{x}^* \\ \mathbf{x}_{k-1} - \mathbf{x}^* \end{bmatrix} = \mathbf{T}^k \begin{bmatrix} \mathbf{x}_1 - \mathbf{x}^* \\ \mathbf{x}_0 - \mathbf{x}^* \end{bmatrix}$$

Proof overview (Polyak HBM)

Let's first look at how to analyze the classical heavy ball momentum algorithm:

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \underbrace{\alpha (\mathbf{A}^\top \mathbf{A} \mathbf{x}_k - \mathbf{A}^\top \mathbf{b})}_{\text{gradient}} + \underbrace{\beta (\mathbf{x}_k - \mathbf{x}_{k-1})}_{\text{momentum}}.$$

Assuming a consistent system, $\mathbf{b} = \mathbf{A} \mathbf{x}^*$, so we can rewrite the update as

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \mathbf{A}^\top \mathbf{A} (\mathbf{x}_k - \mathbf{x}^*) + \beta (\mathbf{x}_k - \mathbf{x}_{k-1}).$$

Therefore,

$$\begin{aligned} \mathbf{x}_{k+1} - \mathbf{x}^* &= (\mathbf{I} - \alpha \mathbf{A}^\top \mathbf{A}) (\mathbf{x}_k - \mathbf{x}^*) + \beta (\mathbf{x}_k - \mathbf{x}^* + \mathbf{x}^* - \mathbf{x}_{k-1}) \\ &= ((1 + \beta) \mathbf{I} - \alpha \mathbf{A}^\top \mathbf{A}) (\mathbf{x}_k - \mathbf{x}^*) - \beta (\mathbf{x}_k - \mathbf{x}_{k-1}). \end{aligned}$$

We can write this as a matrix iteration

$$\begin{bmatrix} \mathbf{x}_{k+1} - \mathbf{x}^* \\ \mathbf{x}_k - \mathbf{x}^* \end{bmatrix} = \underbrace{\begin{bmatrix} (1 + \beta) \mathbf{I} - \alpha \mathbf{A}^\top \mathbf{A} & -\beta \mathbf{I} \\ \mathbf{I} & \mathbf{0} \end{bmatrix}}_{\mathbf{T} = \mathbf{T}(\alpha, \beta)} \begin{bmatrix} \mathbf{x}_k - \mathbf{x}^* \\ \mathbf{x}_{k-1} - \mathbf{x}^* \end{bmatrix} = \mathbf{T}^k \begin{bmatrix} \mathbf{x}_1 - \mathbf{x}^* \\ \mathbf{x}_0 - \mathbf{x}^* \end{bmatrix}$$

Proof overview (Polyak HBM)

Let's first look at how to analyze the classical heavy ball momentum algorithm:

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \underbrace{\alpha (\mathbf{A}^\top \mathbf{A} \mathbf{x}_k - \mathbf{A}^\top \mathbf{b})}_{\text{gradient}} + \underbrace{\beta (\mathbf{x}_k - \mathbf{x}_{k-1})}_{\text{momentum}}.$$

Assuming a consistent system, $\mathbf{b} = \mathbf{A} \mathbf{x}^*$, so we can rewrite the update as

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha \mathbf{A}^\top \mathbf{A} (\mathbf{x}_k - \mathbf{x}^*) + \beta (\mathbf{x}_k - \mathbf{x}_{k-1}).$$

Therefore,

$$\begin{aligned} \mathbf{x}_{k+1} - \mathbf{x}^* &= (\mathbf{I} - \alpha \mathbf{A}^\top \mathbf{A}) (\mathbf{x}_k - \mathbf{x}^*) + \beta (\mathbf{x}_k - \mathbf{x}^* + \mathbf{x}^* - \mathbf{x}_{k-1}) \\ &= ((1 + \beta) \mathbf{I} - \alpha \mathbf{A}^\top \mathbf{A}) (\mathbf{x}_k - \mathbf{x}^*) - \beta (\mathbf{x}_k - \mathbf{x}_{k-1}). \end{aligned}$$

We can write this as a matrix iteration

$$\begin{bmatrix} \mathbf{x}_{k+1} - \mathbf{x}^* \\ \mathbf{x}_k - \mathbf{x}^* \end{bmatrix} = \underbrace{\begin{bmatrix} (1 + \beta) \mathbf{I} - \alpha \mathbf{A}^\top \mathbf{A} & -\beta \mathbf{I} \\ \mathbf{I} & \mathbf{0} \end{bmatrix}}_{\mathbf{T} = \mathbf{T}(\alpha, \beta)} \begin{bmatrix} \mathbf{x}_k - \mathbf{x}^* \\ \mathbf{x}_{k-1} - \mathbf{x}^* \end{bmatrix} = \mathbf{T}^k \begin{bmatrix} \mathbf{x}_1 - \mathbf{x}^* \\ \mathbf{x}_0 - \mathbf{x}^* \end{bmatrix}$$

Proof overview (Polyak HBM)

Now, bounding the error,

$$\left\| \begin{bmatrix} \mathbf{x}_{k+1} - \mathbf{x}^* \\ \mathbf{x}_k - \mathbf{x}^* \end{bmatrix} \right\| \leq \|\mathbf{T}^k\| \left\| \begin{bmatrix} \mathbf{x}_1 - \mathbf{x}^* \\ \mathbf{x}_0 - \mathbf{x}^* \end{bmatrix} \right\|.$$

It's straightforward (but a bit tedious) to show that the eigenvalues of \mathbf{T} are

$$z_j^\pm := \frac{1}{2} \left(1 + \beta - \alpha\lambda_j \pm \sqrt{(1 + \beta - \alpha\lambda_j)^2 - 4\beta} \right).$$

Finally, optimizing the choices of α and β , we can get a bound

$$\|\mathbf{T}^k\| \approx \|\mathbf{T}\|^k = \left(\frac{\sqrt{\lambda_{\max}/\lambda_{\min}} - 1}{\sqrt{\lambda_{\max}/\lambda_{\min}} + 1} \right)^k \leq \exp \left(-\frac{k}{2\sqrt{\lambda_{\max}/\lambda_{\min}}} \right).$$

Proof overview (Polyak HBM)

Now, bounding the error,

$$\left\| \begin{bmatrix} \mathbf{x}_{k+1} - \mathbf{x}^* \\ \mathbf{x}_k - \mathbf{x}^* \end{bmatrix} \right\| \leq \|\mathbf{T}^k\| \left\| \begin{bmatrix} \mathbf{x}_1 - \mathbf{x}^* \\ \mathbf{x}_0 - \mathbf{x}^* \end{bmatrix} \right\|.$$

It's straightforward (but a bit tedious) to show that the eigenvalues of \mathbf{T} are

$$z_j^\pm := \frac{1}{2} \left(1 + \beta - \alpha\lambda_j \pm \sqrt{(1 + \beta - \alpha\lambda_j)^2 - 4\beta} \right).$$

Finally, optimizing the choices of α and β , we can get a bound

$$\|\mathbf{T}^k\| \approx \|\mathbf{T}\|^k = \left(\frac{\sqrt{\lambda_{\max}/\lambda_{\min}} - 1}{\sqrt{\lambda_{\max}/\lambda_{\min}} + 1} \right)^k \leq \exp \left(-\frac{k}{2\sqrt{\lambda_{\max}/\lambda_{\min}}} \right).$$

Proof overview (Polyak HBM)

Now, bounding the error,

$$\left\| \begin{bmatrix} \mathbf{x}_{k+1} - \mathbf{x}^* \\ \mathbf{x}_k - \mathbf{x}^* \end{bmatrix} \right\| \leq \|\mathbf{T}^k\| \left\| \begin{bmatrix} \mathbf{x}_1 - \mathbf{x}^* \\ \mathbf{x}_0 - \mathbf{x}^* \end{bmatrix} \right\|.$$

It's straightforward (but a bit tedious) to show that the eigenvalues of \mathbf{T} are

$$z_j^\pm := \frac{1}{2} \left(1 + \beta - \alpha\lambda_j \pm \sqrt{(1 + \beta - \alpha\lambda_j)^2 - 4\beta} \right).$$

Finally, optimizing the choices of α and β , we can get a bound

$$\|\mathbf{T}^k\| \approx \|\mathbf{T}\|^k = \left(\frac{\sqrt{\lambda_{\max}/\lambda_{\min}} - 1}{\sqrt{\lambda_{\max}/\lambda_{\min}} + 1} \right)^k \leq \exp \left(-\frac{k}{2\sqrt{\lambda_{\max}/\lambda_{\min}}} \right).$$

Proof overview

In minibatch-HBM, we replace the true gradient with a minibatch stochastic gradient

$$\nabla f_{S_k}(\mathbf{x}_k) = \frac{1}{B} \sum_{j \in S_k} \frac{1}{p_j} \mathbf{a}_j \mathbf{a}_j^\top (\mathbf{x}_k - \mathbf{x}^*).$$

Define the random matrix

$$\mathbf{M}_{S_k} = \frac{1}{B} \sum_{j \in S_k} \frac{1}{p_j} \mathbf{a}_j \mathbf{a}_j^\top$$

and observe that

$$\mathbb{E}[\mathbf{M}_{S_k}] = \frac{1}{B} \mathbb{E} \left[\sum_{j \in S_k} \frac{1}{p_j} \mathbf{a}_j \mathbf{a}_j^\top \right] = \mathbf{A}^\top \mathbf{A}.$$

Proof overview

A similar argument gives us a recurrence

$$\begin{bmatrix} \mathbf{x}_{k+1} - \mathbf{x}^* \\ \mathbf{x}_k - \mathbf{x}^* \end{bmatrix} = \underbrace{\begin{bmatrix} (1 + \beta)\mathbf{I} - \alpha\mathbf{M}_{S_k} & -\beta\mathbf{I} \\ \mathbf{I} & \mathbf{0} \end{bmatrix}}_{\mathbf{Y}_{S_k} = \mathbf{Y}_{S_k}(\alpha, \beta)} \begin{bmatrix} \mathbf{x}_k - \mathbf{x}^* \\ \mathbf{x}_{k-1} - \mathbf{x}^* \end{bmatrix}.$$

This gives an error bound

$$\left\| \begin{bmatrix} \mathbf{x}_{k+1} - \mathbf{x}^* \\ \mathbf{x}_k - \mathbf{x}^* \end{bmatrix} \right\| \leq \|\mathbf{Y}_{S_k} \mathbf{Y}_{S_{k-1}} \cdots \mathbf{Y}_{S_1}\| \left\| \begin{bmatrix} \mathbf{x}_1 - \mathbf{x}^* \\ \mathbf{x}_0 - \mathbf{x}^* \end{bmatrix} \right\|.$$

We might hope that

$$\|\mathbf{Y}_{S_k} \mathbf{Y}_{S_{k-1}} \cdots \mathbf{Y}_{S_1}\| \approx \|\mathbb{E}[\mathbf{Y}_{S_k}] \mathbb{E}[\mathbf{Y}_{S_{k-1}}] \cdots \mathbb{E}[\mathbf{Y}_{S_1}]\| = \|\mathbf{T}^k\|.$$

Proof overview

A similar argument gives us a recurrence

$$\begin{bmatrix} \mathbf{x}_{k+1} - \mathbf{x}^* \\ \mathbf{x}_k - \mathbf{x}^* \end{bmatrix} = \underbrace{\begin{bmatrix} (1 + \beta)\mathbf{I} - \alpha\mathbf{M}_{S_k} & -\beta\mathbf{I} \\ \mathbf{I} & \mathbf{0} \end{bmatrix}}_{\mathbf{Y}_{S_k} = \mathbf{Y}_{S_k}(\alpha, \beta)} \begin{bmatrix} \mathbf{x}_k - \mathbf{x}^* \\ \mathbf{x}_{k-1} - \mathbf{x}^* \end{bmatrix}.$$

This gives an error bound

$$\left\| \begin{bmatrix} \mathbf{x}_{k+1} - \mathbf{x}^* \\ \mathbf{x}_k - \mathbf{x}^* \end{bmatrix} \right\| \leq \|\mathbf{Y}_{S_k} \mathbf{Y}_{S_{k-1}} \cdots \mathbf{Y}_{S_1}\| \left\| \begin{bmatrix} \mathbf{x}_1 - \mathbf{x}^* \\ \mathbf{x}_0 - \mathbf{x}^* \end{bmatrix} \right\|.$$

We might hope that

$$\|\mathbf{Y}_{S_k} \mathbf{Y}_{S_{k-1}} \cdots \mathbf{Y}_{S_1}\| \approx \|\mathbb{E}[\mathbf{Y}_{S_k}] \mathbb{E}[\mathbf{Y}_{S_{k-1}}] \cdots \mathbb{E}[\mathbf{Y}_{S_1}]\| = \|\mathbf{T}^k\|.$$

Proof overview

A similar argument gives us a recurrence

$$\begin{bmatrix} \mathbf{x}_{k+1} - \mathbf{x}^* \\ \mathbf{x}_k - \mathbf{x}^* \end{bmatrix} = \underbrace{\begin{bmatrix} (1 + \beta)\mathbf{I} - \alpha\mathbf{M}_{S_k} & -\beta\mathbf{I} \\ \mathbf{I} & \mathbf{0} \end{bmatrix}}_{\mathbf{Y}_{S_k} = \mathbf{Y}_{S_k}(\alpha, \beta)} \begin{bmatrix} \mathbf{x}_k - \mathbf{x}^* \\ \mathbf{x}_{k-1} - \mathbf{x}^* \end{bmatrix}.$$

This gives an error bound

$$\left\| \begin{bmatrix} \mathbf{x}_{k+1} - \mathbf{x}^* \\ \mathbf{x}_k - \mathbf{x}^* \end{bmatrix} \right\| \leq \|\mathbf{Y}_{S_k} \mathbf{Y}_{S_{k-1}} \cdots \mathbf{Y}_{S_1}\| \left\| \begin{bmatrix} \mathbf{x}_1 - \mathbf{x}^* \\ \mathbf{x}_0 - \mathbf{x}^* \end{bmatrix} \right\|.$$

We might hope that

$$\|\mathbf{Y}_{S_k} \mathbf{Y}_{S_{k-1}} \cdots \mathbf{Y}_{S_1}\| \approx \|\mathbb{E}[\mathbf{Y}_{S_k}] \mathbb{E}[\mathbf{Y}_{S_{k-1}}] \cdots \mathbb{E}[\mathbf{Y}_{S_1}]\| = \|\mathbf{T}^k\|.$$

Products of random matrices

Theorem (Huang, Niles-Weed, Tropp, and Ward 2021). Consider an independent sequence of $d \times d$ random matrices $\mathbf{X}_1, \dots, \mathbf{X}_k$, and form the product

$$\mathbf{Z} = \mathbf{X}_k \mathbf{X}_{k-1} \cdots \mathbf{X}_1.$$

Assume $\|\mathbb{E}[\mathbf{X}_i]\| \leq q_i$ and $\mathbb{E}[\|\mathbf{X}_i - \mathbb{E}\mathbf{X}_i\|^2]^{1/2} \leq \sigma_i q_i$ for $i = 1, \dots, k$. Let $Q = \prod_{i=1}^n q_i$ and $v = \sum_{i=1}^k \sigma_i^2$. Then

$$\mathbb{E}[\|\mathbf{Z}\|] \leq Q \exp\left(\sqrt{2v \max\{2v, \log(d)\}}\right).$$

To determine q_i and σ_i we use **matrix concentration** bounds for sums of matrices⁵.

⁵Tropp 2015.

Outlook

Lots of interesting potential follow up work:

- adaptive/automatic parameter selection
- beyond least squares

Key question:

What can we learn about matrices from matvecs?

The matrix-vector product query model

In this **model of computation**, we are only given access to \mathbf{A} via a black box that lets us perform matrix-vector products $\mathbf{x} \mapsto \mathbf{A}\mathbf{x}$ and $\mathbf{x} \mapsto \mathbf{A}^T\mathbf{x}$.

Examples:

- $\mathbf{A} = \mathbf{B}^{-1}$, and we only know \mathbf{B}
 - perform products by solving $\mathbf{B}\mathbf{y} = \mathbf{x}$
 - this generalizes to $\mathbf{A} = f(\mathbf{B})$
- $\mathbf{A} = \exp(-it\mathbf{H})$
 - perform products by evolving according to Hamiltonian \mathbf{H} for time t from initial state \mathbf{x}
- \mathbf{A} is the system matrix for a CT scanner
 - perform products by running the CT scanner on target \mathbf{x}
- \mathbf{A} is a regular matrix
 - perform products with highly optimized hardware/software (e.g. GPUs, crossbar array, MKL, etc.)

The matrix-vector product query model

In this **model of computation**, we are only given access to \mathbf{A} via a black box that lets us perform matrix-vector products $\mathbf{x} \mapsto \mathbf{A}\mathbf{x}$ and $\mathbf{x} \mapsto \mathbf{A}^T\mathbf{x}$.

Examples:

- $\mathbf{A} = \mathbf{B}^{-1}$, and we only know \mathbf{B}
 - perform products by solving $\mathbf{B}\mathbf{y} = \mathbf{x}$
 - this generalizes to $\mathbf{A} = f(\mathbf{B})$
- $\mathbf{A} = \exp(-it\mathbf{H})$
 - perform products by evolving according to Hamiltonian \mathbf{H} for time t from initial state \mathbf{x}
- \mathbf{A} is the system matrix for a CT scanner
 - perform products by running the CT scanner on target \mathbf{x}
- \mathbf{A} is a regular matrix
 - perform products with highly optimized hardware/software (e.g. GPUs, crossbar array, MKL, etc.)

The matrix-vector product query model

In this **model of computation**, we are only given access to \mathbf{A} via a black box that lets us perform matrix-vector products $\mathbf{x} \mapsto \mathbf{A}\mathbf{x}$ and $\mathbf{x} \mapsto \mathbf{A}^T\mathbf{x}$.

Examples:

- $\mathbf{A} = \mathbf{B}^{-1}$, and we only know \mathbf{B}
 - perform products by solving $\mathbf{B}\mathbf{y} = \mathbf{x}$
 - this generalizes to $\mathbf{A} = f(\mathbf{B})$
- $\mathbf{A} = \exp(-it\mathbf{H})$
 - perform products by evolving according to Hamiltonian \mathbf{H} for time t from initial state \mathbf{x}
- \mathbf{A} is the system matrix for a CT scanner
 - perform products by running the CT scanner on target \mathbf{x}
- \mathbf{A} is a regular matrix
 - perform products with highly optimized hardware/software (e.g. GPUs, crossbar array, MKL, etc.)

The matrix-vector product query model

In this **model of computation**, we are only given access to \mathbf{A} via a black box that lets us perform matrix-vector products $\mathbf{x} \mapsto \mathbf{A}\mathbf{x}$ and $\mathbf{x} \mapsto \mathbf{A}^\top \mathbf{x}$.

Examples:

- $\mathbf{A} = \mathbf{B}^{-1}$, and we only know \mathbf{B}
 - perform products by solving $\mathbf{B}\mathbf{y} = \mathbf{x}$
 - this generalizes to $\mathbf{A} = f(\mathbf{B})$
- $\mathbf{A} = \exp(-it\mathbf{H})$
 - perform products by evolving according to Hamiltonian \mathbf{H} for time t from initial state \mathbf{x}
- \mathbf{A} is the system matrix for a CT scanner
 - perform products by running the CT scanner on target \mathbf{x}
- \mathbf{A} is a regular matrix
 - perform products with highly optimized hardware/software (e.g. GPUs, crossbar array, MKL, etc.)

The matrix-vector product query model

In this **model of computation**, we are only given access to \mathbf{A} via a black box that lets us perform matrix-vector products $\mathbf{x} \mapsto \mathbf{A}\mathbf{x}$ and $\mathbf{x} \mapsto \mathbf{A}^\top \mathbf{x}$.

Examples:

- $\mathbf{A} = \mathbf{B}^{-1}$, and we only know \mathbf{B}
 - perform products by solving $\mathbf{B}\mathbf{y} = \mathbf{x}$
 - this generalizes to $\mathbf{A} = f(\mathbf{B})$
- $\mathbf{A} = \exp(-it\mathbf{H})$
 - perform products by evolving according to Hamiltonian \mathbf{H} for time t from initial state \mathbf{x}
- \mathbf{A} is the system matrix for a CT scanner
 - perform products by running the CT scanner on target \mathbf{x}
- \mathbf{A} is a regular matrix
 - perform products with highly optimized hardware/software (e.g. GPUs, crossbar array, MKL, etc.)

The computational complexity of linear algebra

Numerical linear algebra is one of the oldest fields of computing.

- Traditionally, the cost of algorithms is studied in terms of floating point operations, but flops are increasingly less relevant.
- Now, things like matrix products (or matrix loads) often dominate the cost of an algorithm.

If we take matrix-products as the basic unit of computation, we can also study the **computational complexity** of linear algebra problems.

- almost all linear algebra problems can be solved in $O(n^3)$ time, so classical complexity theory (e.g. P vs NP) is not very useful

Structured matrices

Often, we can solve linear algebra problems faster when the matrices of interest are **structured**.

Examples:

- low-rank matrices
- hierarchical matrices
- sparse matrices
- banded matrices

Common framework: Approximate \mathbf{A} with a structured matrix, then use the structured approximation.

Matrix approximation

Problem. Let S be some family of matrices and \mathbf{A} an arbitrary matrix that can only be accessed by matrix-vector products.

Find a matrix $\tilde{\mathbf{A}} \in S$ such that

$$\|\mathbf{A} - \tilde{\mathbf{A}}\| \leq (1 + \varepsilon) \min_{\mathbf{X} \in S} \|\mathbf{A} - \mathbf{X}\|.$$

This problem asks us to find a structured matrix approximation $\tilde{\mathbf{A}}$ to \mathbf{A} competitive with the **best possible** approximation.

- If $\mathbf{A} \in S$, then we require $\tilde{\mathbf{A}} = \mathbf{A}$
- If \mathbf{A} has a good approximation from S , then it might be okay for ε to be large

Wxample: low-rank approximation⁶

⁶Halko, Martinsson, and Tropp 2011; Musco and Musco 2015; Tropp and Webber 2023.

Hierarchical matrix motivation: kernel matrices

Consider data $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$ and a kernel function $k(\mathbf{x}_i, \mathbf{x}_j)$.

- examples: $k(\mathbf{x}, \mathbf{y}) = 1/\|\mathbf{x} - \mathbf{y}\|^2$, $k(\mathbf{x}, \mathbf{y}) = \exp(-\|\mathbf{x} - \mathbf{y}\|^2/\sigma^2)$, $k(\mathbf{x}, \mathbf{y}) = (\mathbf{x}^\top \mathbf{y} + c)^q$

We can define a **kernel matrix**


$$\mathbf{K} = \begin{bmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & \cdots & k(\mathbf{x}_1, \mathbf{x}_n) \\ \vdots & \ddots & \vdots \\ k(\mathbf{x}_n, \mathbf{x}_1) & \cdots & k(\mathbf{x}_n, \mathbf{x}_n) \end{bmatrix}.$$

Kernel matrices are widely used in data science and machine learning applications for task such as clustering and classification.⁷

⁷?

Hierarchical matrix motivation: kernel matrices


For simplicity, consider 1 dimensional data:



Observations:

- Interactions between points in non-adjacent groups are approximately **low-rank**.
- We can **recursively** treat interactions between adjacent groups or within a group at a finer scale.

Hierarchical matrices




low-rank block




recursive block


Hierarchical matrices



Hierarchical matrices



HODLR matrices



The HODLR approximation problem

Problem. Given an $n \times n$ matrix \mathbf{A} , accessible only by matrix-vector products, a rank parameter k , and an accuracy parameter ε , find a $\text{HODLR}(k)$ matrix $\tilde{\mathbf{A}}$ such that

$$\mathbb{E}[\|\mathbf{A} - \tilde{\mathbf{A}}\|_{\text{F}}] \leq (1 + \varepsilon) \min_{\mathbf{H} \in \text{HODLR}(k)} \|\mathbf{A} - \mathbf{H}\|_{\text{F}}.$$

The best HODLR approximation to \mathbf{A} is obtained by applying a rank- k SVD to each low-rank block of \mathbf{A} .

- This is too expensive in the matrix-vector product model (n products)

In the special case that $\mathbf{A} \in \text{HODLR}(k)$, then we require $\tilde{\mathbf{A}} = \mathbf{A}$ (regardless of ε).

- There are several matvec algorithms for this setting⁸

⁸Lin, Lu, and Ying 2011; Martinsson 2016; Levitt and Martinsson 2022; Halikias and Townsend 2023.

The HODLR approximation problem

Problem. Given an $n \times n$ matrix \mathbf{A} , accessible only by matrix-vector products, a rank parameter k , and an accuracy parameter ε , find a $\text{HODLR}(k)$ matrix $\tilde{\mathbf{A}}$ such that

$$\mathbb{E}[\|\mathbf{A} - \tilde{\mathbf{A}}\|_F] \leq (1 + \varepsilon) \min_{\mathbf{H} \in \text{HODLR}(k)} \|\mathbf{A} - \mathbf{H}\|_F.$$

The best HODLR approximation to \mathbf{A} is obtained by applying a rank- k SVD to each low-rank block of \mathbf{A} .

- This is too expensive in the matrix-vector product model (n products)

In the special case that $\mathbf{A} \in \text{HODLR}(k)$, then we require $\tilde{\mathbf{A}} = \mathbf{A}$ (regardless of ε).

- There are several matvec algorithms for this setting⁸

⁸Lin, Lu, and Ying 2011; Martinsson 2016; Levitt and Martinsson 2022; Halikias and Townsend 2023.

The HODLR approximation problem

Problem. Given an $n \times n$ matrix \mathbf{A} , accessible only by matrix-vector products, a rank parameter k , and an accuracy parameter ε , find a $\text{HODLR}(k)$ matrix $\tilde{\mathbf{A}}$ such that

$$\mathbb{E}[\|\mathbf{A} - \tilde{\mathbf{A}}\|_F] \leq (1 + \varepsilon) \min_{\mathbf{H} \in \text{HODLR}(k)} \|\mathbf{A} - \mathbf{H}\|_F.$$

The best HODLR approximation to \mathbf{A} is obtained by applying a rank- k SVD to each low-rank block of \mathbf{A} .

- This is too expensive in the matrix-vector product model (n products)

In the special case that $\mathbf{A} \in \text{HODLR}(k)$, then we require $\tilde{\mathbf{A}} = \mathbf{A}$ (regardless of ε).

- There are several matvec algorithms for this setting⁸

⁸Lin, Lu, and Ying 2011; Martinsson 2016; Levitt and Martinsson 2022; Halikias and Townsend 2023.

Learning low-rank matrices from matrix-vector products

The Randomized SVD (RSVD) is a well-known algorithm for obtaining a low-rank approximation to a matrix \mathbf{B} :

1. Sample Gaussian matrix $\mathbf{\Omega}$
2. Form $\mathbf{Q} = \text{orth}(\mathbf{B}\mathbf{\Omega})$
3. Compute $\mathbf{X} = \mathbf{Q}^\top \mathbf{B}$
4. Output $\mathbf{Q}\llbracket \mathbf{X} \rrbracket_k$

Theorem. If $\mathbf{\Omega}$ has $O(k/\varepsilon)$ columns, then

$$\|\mathbf{B} - \mathbf{Q}\llbracket \mathbf{X} \rrbracket_k\|_F \leq (1 + \varepsilon) \min_{\text{rank}(\mathbf{X}) \leq k} \|\mathbf{B} - \mathbf{X}\|_F.$$

Corollary. If \mathbf{B} is rank- k , then $\mathbf{Q}\llbracket \mathbf{X} \rrbracket_k = \mathbf{B}$ (with probability one).

Peeling: an algorithm for the recovery problem⁹


The algorithm works from the top layer down.

At each level, we **simultaneously** apply the RSVD to the low-rank off-diagonal blocks.

We then “peel” off these blocks before proceeding to the next level

⁹Lin, Lu, and Ying 2011; Martinsson 2016.

Peeling: an algorithm for the recovery problem



Peeling: an algorithm for the recovery problem

At each level we use k matrix-vector products with \mathbf{A} and \mathbf{A}^\top .


There are $\log_2(n/k) \leq \log_2(n)$ levels until the blocks are of size k

- then we can directly recover them at once with k products

Theorem (Lin, Lu, and Ying 2011). There is an algorithm to recover a HODLR matrix using $O(k \log_2(n))$ products with \mathbf{A} .

If \mathbf{A} is not HODLR, we can still apply this algorithm, but we might worry about **exponential blow-up** in the error!

Peeling with error



We obtain perturbed sketches of the form $\mathbf{A}_{6,5}^{(3)}\Omega_5 + \mathbf{A}_{6,1}^{(3)}\Omega_1 + \mathbf{A}_{6,3}^{(3)}\Omega_3 + \mathbf{A}_{6,7}^{(3)}\Omega_7$ and $(\mathbf{A}_{6,5}^{(3)})^\top\Psi_6 + (\mathbf{A}_{2,5}^{(3)})^\top\Psi_2 + (\mathbf{A}_{4,5}^{(3)})^\top\Psi_4 + (\mathbf{A}_{8,5}^{(3)})^\top\Psi_8$.

The RSVD tries to compute $\mathbf{Q}^\top \mathbf{B}$ directly; this is the solution to:

$$\min_{\mathbf{X}} \|\mathbf{A} - \mathbf{Q}\mathbf{X}\|_F.$$

Instead, we can solve a sketched problem:

$$\min_{\mathbf{X}} \|\Psi^\top \mathbf{A} - \Psi^\top \mathbf{Q}\mathbf{X}\|_F.$$

This means $\mathbf{X} = (\Psi^\top \mathbf{Q})^\dagger \Psi^\top \mathbf{A}$.

Observation. By adding columns to Ψ , we can average out certain errors in the product $\Psi^\top \mathbf{A}$.

The algorithm is also non-adaptive (we can do products with Ψ in advance)

¹⁰Clarkson and Woodruff 2009; Tropp, Yurtsever, Udell, and Cevher 2017; Nakatsukasa 2020.

The RSVD tries to compute $\mathbf{Q}^\top \mathbf{B}$ directly; this is the solution to:

$$\min_{\mathbf{X}} \|\mathbf{A} - \mathbf{Q}\mathbf{X}\|_F.$$

Instead, we can solve a sketched problem:

$$\min_{\mathbf{X}} \|\Psi^\top \mathbf{A} - \Psi^\top \mathbf{Q}\mathbf{X}\|_F.$$

This means $\mathbf{X} = (\Psi^\top \mathbf{Q})^\dagger \Psi^\top \mathbf{A}$.

Observation. By adding columns to Ψ , we can average out certain errors in the product $\Psi^\top \mathbf{A}$.

The algorithm is also non-adaptive (we can do products with Ψ in advance)

¹⁰Clarkson and Woodruff 2009; Tropp, Yurtsever, Udell, and Cevher 2017; Nakatsukasa 2020.

The RSVD tries to compute $\mathbf{Q}^\top \mathbf{B}$ directly; this is the solution to:

$$\min_{\mathbf{X}} \|\mathbf{A} - \mathbf{Q}\mathbf{X}\|_F.$$

Instead, we can solve a sketched problem:

$$\min_{\mathbf{X}} \|\Psi^\top \mathbf{A} - \Psi^\top \mathbf{Q}\mathbf{X}\|_F.$$

This means $\mathbf{X} = (\Psi^\top \mathbf{Q})^\dagger \Psi^\top \mathbf{A}$.

Observation. By adding columns to Ψ , we can average out certain errors in the product $\Psi^\top \mathbf{A}$.

The algorithm is also non-adaptive (we can do products with Ψ in advance)

¹⁰Clarkson and Woodruff 2009; Tropp, Yurtsever, Udell, and Cevher 2017; Nakatsukasa 2020.

A perturbation bound for the RSVD

We prove a perturbation bound for the RSVD. This is likely of independent interest.

Theorem (Chen, Keles, Halikias, Musco, Musco, and Persson 2024). Let $\mathbf{Q} = \text{orth}(\mathbf{B}\mathbf{\Omega} + \mathbf{E}_1)$ and $\mathbf{X} = \mathbf{Q}^\top \mathbf{B} + \mathbf{E}_2$. Then

$$\|\mathbf{B} - \mathbf{Q}[\mathbf{Q}^\top \mathbf{B} + \mathbf{E}_2]_k\|_F \leq \underbrace{\|\mathbf{E}_1 \mathbf{\Omega}_{\text{top}}^\dagger\|_F + 2\|\mathbf{E}_2\|_F}_{\text{perturbations}} + \underbrace{\left(\|\mathbf{\Sigma}_{\text{bot}}\|_F^2 + \|\mathbf{\Sigma}_{\text{bot}} \mathbf{\Omega}_{\text{bot}} \mathbf{\Omega}_{\text{top}}^\dagger\|_F^2\right)^{1/2}}_{\text{classical RSVD bound}}.$$

Takeaway: The pseudoinverse will help damp the perturbation \mathbf{E}_1 , but (unsurprisingly) all of the perturbation \mathbf{E}_2 can propagate.

Matrix approximation


Theorem (Chen, Keles, Halikias, Musco, Musco, and Persson 2024). There exist matvec algorithms which use $O(k \log(n) \cdot \text{poly}(1/\beta))$ products with \mathbf{A} to obtain a HODLR(k) matrix $\widetilde{\mathbf{A}}$ satisfying

$$\|\mathbf{A} - \widetilde{\mathbf{A}}\|_{\text{F}} \leq (1 + \beta)^{\log_2(n)} \min_{\mathbf{H} \in \text{HODLR}(k)} \|\mathbf{A} - \mathbf{H}\|_{\text{F}}.$$

Corollary. $(1+\varepsilon)$ -optimal approximation with $O(k \log(n) \cdot \text{poly}(\log(n)/\varepsilon))$ matvecs

Corollary. $n^{0.01}$ -optimal approximation with $O(k \log(n))$ matvecs

Example



solid = Nyström based algorithm, dashed = RSVD based algorithm

Analysis of peeling

- At each level, we use our perturbation bound to show that we obtain a near-optimal approximation to the blocks at that level, with additional error depending on the quality of approximation at previous levels
- We then analyze the propagation of error through the levels
- The error at any one given level is bounded in terms of the overall best possible error (truncated SVD on each low-rank block)

The fixed-sparsity approximation problem

Fix a $n \times n$ binary matrix \mathbf{S} , and consider the set of matrices whose nonzero pattern is given by \mathbf{S} ; i.e. $\mathcal{S} = \{\mathbf{X} : \mathbf{X} \circ \mathbf{S} = \mathbf{X}\}$, where “ \circ ” is entrywise product.

- examples: diagonal matrices, banded matrices

Problem. Given an $n \times n$ matrix \mathbf{A} , accessible only by matrix-vector products, find a matrix $\tilde{\mathbf{A}}$ with sparsity \mathbf{S} such that

$$\|\mathbf{A} - \tilde{\mathbf{A}}\|_F \leq (1 + \varepsilon) \min_{\mathbf{X} \circ \mathbf{S} = \mathbf{X}} \|\mathbf{A} - \mathbf{X}\|_F.$$

The best sparse approximation to \mathbf{A} is $\mathbf{A} \circ \mathbf{S}$.

- This is too expensive in the matrix-vector product model (n products)

Lots of past work for exact recovery¹¹ and diagonal estimation¹²

¹¹Coleman and Moré 1983; Coleman and Moré 1983.

¹²Baston and Nakatsukasa 2022; Hallman, Ipsen, and Saibaba 2023; Dharangutte and Musco 2023.

The fixed-sparsity approximation problem: an algorithm

One might try a simple **sketching** algorithm:

$$\widetilde{\mathbf{A}} = \operatorname{argmin}_{\mathbf{X} \circ \mathbf{S} = \mathbf{X}} \|\mathbf{A}\mathbf{G} - \mathbf{X}\mathbf{G}\|_{\mathrm{F}}.$$

Theorem (Amsel, Chen, Keles, Halikias, Musco, and Musco 2024). Suppose \mathbf{S} has at most s nonzeros per row. There exists a matvec algorithms which use $O(s/\varepsilon)$ products with \mathbf{A} to obtain a matrix $\widetilde{\mathbf{A}}$ with sparsity \mathbf{S} satisfying

$$\mathbb{E}[\|\mathbf{A} - \widetilde{\mathbf{A}}\|_{\mathrm{F}}] \leq (1 + \varepsilon) \min_{\mathbf{X} \circ \mathbf{S} = \mathbf{X}} \|\mathbf{A} - \mathbf{X}\|_{\mathrm{F}}.$$

Unlike compressed-sensing type problems, **no dimension dependence!**

- in the compressed sensing setting, we have a necessary $\log(n)$ factor

The fixed-sparsity approximation problem (hardness)

We also prove a **lower-bound**, which reveals that our algorithm is optimal, up to constant factors.

Theorem (Amsel, Chen, Keles, Halikias, Musco, and Musco 2024). For any \mathbf{S} with $\Theta(s)$ nonzeros per row/column, there exist hard instances for which **any algorithm** producing $\tilde{\mathbf{A}}$ with sparsity \mathbf{S} must use $\Omega(s/\varepsilon)$ products with \mathbf{A} in order to guarantee

$$\mathbb{P}\left[\|\mathbf{A} - \tilde{\mathbf{A}}\|_{\text{F}} \leq (1 + \varepsilon) \min_{\mathbf{X} \circ \mathbf{S} = \mathbf{X}} \|\mathbf{A} - \mathbf{X}\|_{\text{F}}\right] > 1/2.$$

Lower bound case study: diagonal approximation

To illustrate the key ideas, let's consider the special case $\mathbf{S} = \mathbf{I}$; i.e. diagonal approximation.

Theorem. There exists a constant $C > 0$ such that the following holds:

For any $\varepsilon > 0$, there exists a distribution \mathbf{A} on $d \times d$ matrices such that, any algorithm which uses $m < C/\varepsilon$ queries cannot output an approximation \mathbf{d} to $\text{diag}(\mathbf{A})$ satisfying

$$\mathbb{P}[\|\text{diag}(\mathbf{A}) - \mathbf{d}\|_2^2 \leq \varepsilon \|\mathbf{A} - \mathbf{A} \circ \mathbf{I}\|_F^2] > 1/2.$$

If $\widetilde{\mathbf{A}} = \text{diag}(\mathbf{d})$ and the event in the probability holds,

$$\|\mathbf{A} - \widetilde{\mathbf{A}}\|_F^2 = \|\mathbf{A} \circ \mathbf{S} - \widetilde{\mathbf{A}}\|_F^2 + \|\mathbf{A} - \mathbf{A} \circ \mathbf{I}\|_F^2 \leq (1 + \varepsilon) \|\mathbf{A} - \mathbf{A} \circ \mathbf{I}\|_F^2.$$

Since $\sqrt{1 + \varepsilon} = 1 + O(\varepsilon)$, this is also a lower bound for the approximation problem.

Gaussian matrices after adaptive queries

Lemma. Let $\mathbf{A} \sim \text{Gaussian}(d, d)$. Suppose we make a sequence of adaptive queries $\mathbf{x}_1, \dots, \mathbf{x}_m$ with responses $\mathbf{y}_1 = \mathbf{A}\mathbf{x}_1, \dots, \mathbf{y}_m = \mathbf{A}\mathbf{x}_m$.

Let $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_m]$ and $\mathbf{Y} = [\mathbf{y}_1, \dots, \mathbf{y}_m]$. Then, \mathbf{A} can be factored as

$$\mathbf{A} = [\mathbf{Y} \quad \mathbf{G}] \begin{bmatrix} \mathbf{X}^\top \\ \mathbf{Z}^\top \end{bmatrix},$$

where $\mathbf{Z}^\top \mathbf{X} = \mathbf{0}$, $\mathbf{Z}^\top \mathbf{Z} = \mathbf{I}$, and $\mathbf{G} \sim \text{Gaussian}(d, d - m)$, independently of \mathbf{X} and \mathbf{Y} .

Conditioned on the queries and responses, \mathbf{A} still has a lot of randomness!

Analysis

We have that

$$\mathbf{A} = \mathbf{Y}\mathbf{X}^\top + \mathbf{G}\mathbf{Z}^\top,$$

where $\mathbf{Z}^\top\mathbf{X} = \mathbf{0}$, $\mathbf{Z}^\top\mathbf{Z} = \mathbf{I}$, and $\mathbf{G} \sim \text{Gaussian}(d, d - m)$, independently of \mathbf{X} and \mathbf{Y} .

Conditioned on the queries and responses,

$$\mathbf{A}_{i,i} = \mathbf{e}_i^\top \mathbf{Y}\mathbf{X}^\top \mathbf{e}_i + \mathbf{e}_i^\top \mathbf{G}\mathbf{Z}^\top \mathbf{e}_i = \text{deterministic} + \mathbf{g}_i^\top \mathbf{z}_i.$$

Observe that the \mathbf{g}_i are all independent Gaussian vectors of length $d - m$. Hence

$$d_i = \mathbf{g}_i^\top \mathbf{z}_i \sim N(0, \|\mathbf{z}_i\|^2) \text{ independently.}$$

Suppose we try to output an approximation \mathbf{d} to $\text{diag}(\mathbf{A})$. Based on what we know (queries and responses), the best thing we can do is output $\mathbf{d} = \text{diag}(\mathbf{Y}\mathbf{X}^\top)$. So then

$$\mathbb{E}[\|\text{diag}(\mathbf{A}) - \mathbf{d}\|_2^2] = \sum_{i=1}^n \mathbb{E}[d_i^2] = \sum_{i=1}^n \|\mathbf{z}_i\|^2 = \|\mathbf{Z}\|_F^2 = d - m.$$

Choosing the parameters

Applying concentration inequality,

$$\mathbb{P}\left[\|\text{diag}(\mathbf{A}) - \mathbf{d}\|_2^2 > 100(d - m)\right] \leq \text{small}.$$

In addition,

$$\mathbb{P}\left[\|\mathbf{A} - \mathbf{A} \circ \mathbf{I}\|_{\text{F}}^2 < d^2/100\right] \leq \text{small}.$$

If we set $d = 5000/\varepsilon$ and $m < 2500/\varepsilon$,

$$100(d - m) > 100d/2 = \varepsilon d^2/100.$$

Therefore, by a union bound,

$$\mathbb{P}\left[\|\text{diag}(\mathbf{A}) - \mathbf{d}\|_2^2 < \varepsilon \|\mathbf{A} - \mathbf{A} \circ \mathbf{I}\|_{\text{F}}^2\right] \leq \text{small}.$$

Conclusion

We saw two broad directions in randomized linear algebra that I have worked on.

I'm also interested in many other problems including (but certainly not limited to):

- (partial) trace approximation
- spectrum approximation
- classical Krylov subspace methods

References I

- Amsel, Noah et al. (2024). *Fixed-sparsity matrix approximation from matrix-vector products*.
- Baston, Robert A. and Yuji Nakatsukasa (2022). *Stochastic diagonal estimation: probabilistic bounds and an improved algorithm*.
- Bollapragada, Raghu, Tyler Chen, and Rachel Ward (2024). “On the fast convergence of minibatch heavy ball momentum”. In: *IMA Journal of Numerical Analysis*. To appear.
- Braverman, Vladimir, Aditya Krishnan, and Christopher Musco (June 2022). “Sublinear time spectral density estimation”. In: *Proceedings of the 54th Annual ACM SIGACT Symposium on Theory of Computing*. arXiv cs.DS 2104.03461. ACM.
- Campisi, Michele, David Zueco, and Peter Talkner (Oct. 2010). “Thermodynamic anomalies in open quantum systems: Strong coupling effects in the isotropic XY model”. In: *Chemical Physics* 375.2–3, pp. 187–194.
- Chen, Tyler and Yu-Chen Cheng (Aug. 2022). “Numerical computation of the equilibrium-reduced density matrix for strongly coupled open quantum systems”. In: *The Journal of Chemical Physics* 157.6, p. 064106.
- Chen, Tyler, Thomas Trogon, and Shashanka Ubaru (July 2021). “Analysis of stochastic Lanczos quadrature for spectrum approximation”. In: *Proceedings of the 38th International Conference on Machine Learning*. Vol. 139. Proceedings of Machine Learning Research. PMLR, pp. 1728–1739.
- (2022). *Randomized matrix-free quadrature for spectrum and spectral sum approximation*.
- Chen, Tyler et al. (2024). *Near-optimal hierarchical matrix approximation from matrix-vector products*.
- Clarkson, Kenneth L. and David P. Woodruff (May 2009). “Numerical linear algebra in the streaming model”. In: *Proceedings of the forty-first annual ACM symposium on Theory of computing*. STOC '09. ACM.

References II

- Coleman, Thomas F. and Jorge J. Moré (Feb. 1983). “Estimation of Sparse Jacobian Matrices and Graph Coloring Blems”. In: *SIAM Journal on Numerical Analysis* 20.1, pp. 187–209.
- Dharangutte, Prathamesh and Christopher Musco (Jan. 2023). “A Tight Analysis of Hutchinson’s Diagonal Estimator”. In: *Symposium on Simplicity in Algorithms (SOSA)*. Society for Industrial and Applied Mathematics, pp. 353–364.
- Halikias, Diana and Alex Townsend (Sept. 2023). “Structured matrix recovery from matrix-vector products”. In: *Numerical Linear Algebra with Applications* 31.1.
- Halko, N., P. G. Martinsson, and J. A. Tropp (2011). “Finding structure with randomness: probabilistic algorithms for constructing approximate matrix decompositions”. In: *SIAM Rev.* 53.2, pp. 217–288.
- Hallman, Eric, Ilse C. F. Ipsen, and Arvind K. Saibaba (Mar. 2023). “Monte Carlo Methods for Estimating the Diagonal of a Real Symmetric Matrix”. In: *SIAM Journal on Matrix Analysis and Applications* 44.1, pp. 240–269.
- Han, Insu et al. (Jan. 2017). “Approximating Spectral Sums of Large-Scale Matrices using Stochastic Chebyshev Approximations”. In: *SIAM Journal on Scientific Computing* 39.4, A1558–A1585.
- Huang, De et al. (2021). “Matrix concentration for products”. In: *Foundations of Computational Mathematics*, pp. 1–33.
- Ingold, Gert-Ludwig, Peter Hänggi, and Peter Talkner (June 2009). “Specific heat anomalies of open quantum systems”. In: *Physical Review E* 79.6.
- Jain, Prateek et al. (2018). “Accelerating stochastic gradient descent for least squares regression”. In: *Conference On Learning Theory*. PMLR, pp. 545–604.
- Levitt, James and Per-Gunnar Martinsson (2022). *Randomized Compression of Rank-Structured Matrices Accelerated with Graph Coloring*.

References III

- Li, Ruipeng et al. (Jan. 2019). “The Eigenvalues Slicing Library (EVSU): Algorithms, Implementation, and Software”. In: *SIAM Journal on Scientific Computing* 41.4, pp. C393–C415.
- Lin, Lin, Jianfeng Lu, and Lexing Ying (May 2011). “Fast construction of hierarchical matrix representation from matrix–vector multiplication”. In: *Journal of Computational Physics* 230.10, pp. 4071–4087.
- Martinsson, Per-Gunnar (Jan. 2016). “Compressing Rank-Structured Matrices via Randomized Sampling”. In: *SIAM Journal on Scientific Computing* 38.4, A1959–A1986.
- Musco, Cameron and Christopher Musco (2015). *Randomized Block Krylov Methods for Stronger and Faster Approximate Singular Value Decomposition*. Montreal, Canada.
- Nakatsukasa, Yuji (2020). “Fast and stable randomized low-rank matrix approximation”. In: *ArXiv* abs/2009.11392.
- Polizzi, Eric (Mar. 2009). “Density-matrix-based algorithm for solving eigenvalue problems”. In: *Physical Review B* 79.11.
- Talkner, Peter and Peter Hänggi (Oct. 2020). “Colloquium : Statistical mechanics and thermodynamics at strong coupling: Quantum and classical”. In: *Reviews of Modern Physics* 92.4.
- Tropp, Joel A. (2015). “An Introduction to Matrix Concentration Inequalities”. In: *Foundations and Trends® in Machine Learning* 8.1-2, pp. 1–230.
- Tropp, Joel A. and Robert J. Webber (2023). *Randomized algorithms for low-rank matrix approximation: Design, analysis, and applications*.
- Tropp, Joel A. et al. (Jan. 2017). “Practical Sketching Algorithms for Low-Rank Matrix Approximation”. In: *SIAM Journal on Matrix Analysis and Applications* 38.4, pp. 1454–1485.

References IV

Ubaru, Shashanka, Jie Chen, and Yousef Saad (Jan. 2017). “Fast Estimation of $\text{tr}(f(A))$ via Stochastic Lanczos Quadrature”. In: *SIAM Journal on Matrix Analysis and Applications* 38.4, pp. 1075–1099.

Quantum equilibrium thermodynamics

Consider a quantum system consisting of subsystems (s) and (b) with Hamiltonian

$$\mathbf{H} = \bar{\mathbf{H}}_s + \bar{\mathbf{H}}_b + \mathbf{H}_{sb}, \quad \bar{\mathbf{H}}_s = \mathbf{H}_s \otimes \mathbf{I}_b, \quad \bar{\mathbf{H}}_b = \mathbf{I}_s \otimes \mathbf{H}_b. \quad (1)$$

In thermal equilibrium at interver temperature β , the state of the system is described by a density matrix

$$\rho_t(\beta) = \frac{\exp(-\beta\mathbf{H})}{Z_t(\beta)}, \quad Z_t(\beta) = \text{tr}(\exp(-\beta\mathbf{H})); \quad (2)$$

The denisty matrix for subsystem (s) is given by

$$\rho^*(\beta) = \text{tr}_b(\rho_t(\beta)) = \frac{\text{tr}_b(\exp(-\beta\mathbf{H}))}{\text{tr}(\exp(-\beta\mathbf{H}))}, \quad (3)$$

where $\text{tr}_b(\cdot)$ is the *partial trace* over subsystem (b).¹³

¹³Campisi, Zueco, and Talkner 2010; Ingold, Hänggi, and Talkner 2009; Talkner and Hänggi 2020.

von Neumann entropy of Heisenberg spin chains

The von Neumann entropy $-\text{tr}(\rho^*(\beta) \ln(\rho^*(\beta)))$ is a measure of the **entanglement** between subsystems (s) and (b).

Understanding the von Neumann entropy for a range of a system with Hamiltonian $\mathbf{H}(\theta)$ at a range of parameter values θ and inverse temperatures β is of interest.

We will consider a special case

$$\mathbf{H} = \sum_{i,j} [J_{i,j}^x \sigma_i^x \sigma_j^x + J_{i,j}^y \sigma_i^y \sigma_j^y + J_{i,j}^z \sigma_i^z \sigma_j^z] + \frac{h}{2} \sum_{i=1}^N \sigma_i^z.$$

where h is the magnetic field strength.

Subsystem (s) corresponds to $i = 1, 2$ and subsystem (b) corresponds to the rest of the spins.

Key question:

How to compute reduced density matrices numerically?

A starting point: stochastic trace estimation

If \mathbf{b} is a standard Gaussian random vector:

$$\mathbb{E}[\mathbf{b}^\top f(\mathbf{A})\mathbf{b}] = \text{tr}(f(\mathbf{A})), \quad \mathbb{V}[\mathbf{b}^\top f(\mathbf{A})\mathbf{b}] = 2\|f(\mathbf{A})\|_{\text{F}}^2.$$

It's standard to use a KSM to approximate products $\mathbf{b} \mapsto \mathbf{b}^\top f(\mathbf{A})\mathbf{b}$.

Lots of work balancing the cost of the KSM with the variance of the estimator¹⁴.

¹⁴Han, Malioutov, Avron, and Shin 2017; Ubaru, Chen, and Saad 2017; Chen, Trogon, and Ubaru 2021; Chen, Trogon, and Ubaru 2022; Braverman, Krishnan, and Musco 2022.

Partial traces

Suppose \mathbf{A} is a $d_s d_b \times d_s d_b$ matrix partitioned as:

$$\mathbf{A} = \begin{bmatrix} \mathbf{A}_{1,1} & \mathbf{A}_{1,2} & \cdots & \mathbf{A}_{1,d_s} \\ \mathbf{A}_{2,1} & \mathbf{A}_{2,2} & \cdots & \mathbf{A}_{2,d_s} \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{A}_{d_s,1} & \mathbf{A}_{d_s,2} & \cdots & \mathbf{A}_{d_s,d_s} \end{bmatrix},$$

Partial traces

Then the partial trace (wrt. this partitioning) is defined as:

$$\mathrm{tr}_b(\mathbf{A}) = \begin{bmatrix} \mathrm{tr}(\mathbf{A}_{1,1}) & \mathrm{tr}(\mathbf{A}_{1,2}) & \cdots & \mathrm{tr}(\mathbf{A}_{1,d_s}) \\ \mathrm{tr}(\mathbf{A}_{2,1}) & \mathrm{tr}(\mathbf{A}_{2,2}) & \cdots & \mathrm{tr}(\mathbf{A}_{2,d_s}) \\ \vdots & \vdots & \ddots & \vdots \\ \mathrm{tr}(\mathbf{A}_{d_s,1}) & \mathrm{tr}(\mathbf{A}_{d_s,2}) & \cdots & \mathrm{tr}(\mathbf{A}_{d_s,d_s}) \end{bmatrix}.$$


We can use a randomized estimator:

$$(\mathbf{I}_{d_s} \otimes \mathbf{b})^\top \mathbf{A} (\mathbf{I}_{d_s} \otimes \mathbf{b}) = \begin{bmatrix} \mathbf{b}^\top \mathbf{A}_{1,1} \mathbf{b} & \mathbf{b}^\top \mathbf{A}_{1,2} \mathbf{b} & \cdots & \mathbf{b}^\top \mathbf{A}_{1,d_s} \mathbf{b} \\ \mathbf{b}^\top \mathbf{A}_{2,1} \mathbf{b} & \mathbf{b}^\top \mathbf{A}_{2,2} \mathbf{b} & \cdots & \mathbf{b}^\top \mathbf{A}_{2,d_s} \mathbf{b} \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{b}^\top \mathbf{A}_{d_s,1} \mathbf{b} & \mathbf{b}^\top \mathbf{A}_{d_s,2} \mathbf{b} & \cdots & \mathbf{b}^\top \mathbf{A}_{d_s,d_s} \mathbf{b} \end{bmatrix}.$$

Then use a KSM to approximate products with $\mathbf{A} = f(\mathbf{H})$.

¹⁵Chen and Cheng 2022.

von Neumann entropy



Spectrum approximation

Recall a symmetric matrix \mathbf{A} has a real spectrum $\{\lambda_i\}$. We can encode the spectrum in the **spectral density** function

$$\varphi(x) = \sum_{i=1}^n \frac{1}{n} \delta(x - \lambda_i).$$

Here $\delta(x)$ is a Dirac delta point mass centered at zero.

Problem. Given only matrix-vector product access to \mathbf{A} , approximate $\varphi(x)$.

Example application: high performance computing

State of the art parallel eigensolvers such as FEAST and EVSL work by splitting the spectrum of \mathbf{A} into pieces, which can each be solved on different machines in parallel.¹⁶




The spectral density tells us how many eigenvalues in a region

$$\text{number of eigenvalues in } [a, b] = n \int_a^b \varphi(x) dx.$$

¹⁶Polizzi 2009; Li, Xi, Erlandson, and Saad 2019.

Example application: high performance computing

State of the art parallel eigensolvers such as FEAST and EVSL work by splitting the spectrum of \mathbf{A} into pieces, which can each be solved on different machines in parallel.¹⁶




The spectral density tells us how many eigenvalues in a region

$$\text{number of eigenvalues in } [a, b] = n \int_a^b \varphi(x) dx.$$

¹⁶Polizzi 2009; Li, Xi, Erlandson, and Saad 2019.

Example application: high performance computing

State of the art parallel eigensolvers such as FEAST and EVSL work by splitting the spectrum of \mathbf{A} into pieces, which can each be solved on different machines in parallel.¹⁶




The spectral density tells us how many eigenvalues in a region

$$\text{number of eigenvalues in } [a, b] = n \int_a^b \varphi(x) dx.$$

¹⁶Polizzi 2009; Li, Xi, Erlandson, and Saad 2019.

Wasserstein distance

To measure the error in approximating $\varphi(x)$, we use the **Wasserstein distance**.



Stochastic Lanczos Quadrature

A widely used algorithm for approximating spectral densities Stochastic Lanczos Quadrature (SLQ).

- used in ML to study the training of neural networks
- used in quantum physics for understanding many different systems

Until recently, no theoretical guarantees!

Theorem (Chen, Trogon, and Ubaru 2021). For large matrices,^a SLQ produces an ε accurate approximation to $\varphi(x)$ using $O(1/\varepsilon)$ products with \mathbf{A} .

^aspecifically, for $n \gg 1/\varepsilon^2$

Proof sketch

We use a general technique which actually applies to some other algorithms such as the Kernel Polynomial Method¹⁷.

- If two distributions have similar **Chebyshev moments** through degree $O(1/\varepsilon)$, they are ε close in Wasserstein distance.
 - This follows from Jackson's Theorem from approximation theory and the dual representation of the Wasserstein distance in terms of 1-Lipshitz functions
- The Chebyshev moments of $\varphi(x)$ are $\int T_k(x)\varphi(x)dx = n^{-1} \text{tr}(T_k(\mathbf{A}))$
- We can estimate $\text{tr}(T_k(\mathbf{A}))$ using stochastic trace estimation
 - If \mathbf{x} is an isotropic random vector, $\mathbb{E}[\mathbf{x}^\top \mathbf{M} \mathbf{x}] = \text{tr}(\mathbf{M})$ and $\mathbb{V}[\mathbf{x}^\top \mathbf{M} \mathbf{x}] \approx 2\|\mathbf{M}\|_F$.
- We can compute $\mathbf{x}^\top T_k(\mathbf{A}) \mathbf{x}$ using $O(k)$ products with \mathbf{A} .

¹⁷Chen, Trogon, and Ubaru 2022.

Chebyshev moments

